

aivancity

School of AI & Data for Business & Society

Paris Region · Nice Côte d'Azur

Grande Ecole Programme PGE5 — AI & Data Science

**Auditing the Calibration of
Eligibility Scores
Produced by Large Language
Models
in Patient–Clinical Trial Matching**

*Assessing the Effect of RLHF and Uncovering Mode
Collapse in Open-Source Medical LLMs*

Author: Régis Likassi

Academic Supervisor: Dr. Anuradha Kar

Institution: aivancity Paris

Academic Year 2025–2026

September 2026

Acknowledgements

I would like to express my sincere gratitude to my academic supervisor, Dr. Anuradha Kar, for her guidance, insightful feedback, and constant support throughout this research project. Her expertise in artificial intelligence and her rigorous methodological approach were invaluable in shaping the direction of this work.

I am equally grateful to Aivancity — School of AI & Data for Business & Society for providing the academic environment and resources that made this project possible, and to the Grande Ecole Programme faculty for the training that prepared me for such a research endeavor.

Abstract

Keywords: large language models, calibration, clinical trial matching, RLHF, medical AI, mode collapse, expected calibration error, conformal prediction

Résumé

Mots-clés : grands modèles de langage, calibration, appariement d’essais cliniques, RLHF, IA médicale, mode collapse, expected calibration error, prédiction conforme

Contents

Acknowledgements	i
Abstract	ii
Résumé	iii
List of Abbreviations	x
1 Introduction	1
1.1 Macro Context: AI in Healthcare	1
1.2 Micro Context: Clinical Trial Matching	1
1.3 Problem Statement	1
1.4 Research Question and Hypotheses	1
1.5 Contributions	1
1.6 Structure of the Thesis	1
2 State of the Art	2
2.1 Patient–Clinical Trial Matching	3
2.1.1 Clinical Context and Need	3
2.1.2 Pre-LLM Approaches	3
2.1.3 LLM-based Approaches	3
2.2 Probabilistic Calibration of Classifiers	3
2.2.1 Definitions and Metrics	3
2.2.2 Clinical Importance of Calibration	3
2.2.3 Recalibration Methods	3
2.3 Calibration of Large Language Models	3
2.3.1 The RLHF Effect on Calibration	3
2.3.2 Calibration of Medical LLMs	3
2.3.3 Identified Gap	3
2.4 Open-Source Medical LLMs	3
2.4.1 The Meditron Family	3
2.4.2 BioMistral and Model Merges	3
2.4.3 Other Open-Source Medical LLMs	3
2.5 Positioning of This Work	3
3 Project Context and Scope	5

3.1	Academic Research Framework	5
3.2	Project Origin and Evolution	5
3.3	Dataset: TrialGPT-Criterion-Annotations	5
3.3.1	Dataset Description	5
3.3.2	Justification of the Single-Dataset Choice	5
3.4	Ethical Considerations	5
3.5	Constraints and Resources	5
3.5.1	Computational Constraints	5
3.5.2	Model Access Constraints	5
3.5.3	Time Constraints	5
3.6	Scope of This Work	5
4	Methodology and Technical Approach	6
4.1	Overview of the Experimental Design	6
4.2	Research Question and Pre-registered Hypotheses	6
4.3	Data Preparation	7
4.4	Model Selection	7
4.4.1	Selection Criteria	7
4.4.2	The Six Selected Models	7
4.4.3	Excluded Models	7
4.5	Prompt Design	7
4.5.1	Rationale for Two Prompt Variants	7
4.5.2	Prompt A: Generic Baseline	7
4.5.3	Prompt B: Adapted with Expert Guidance	7
4.5.4	Adaptations for Base vs Instruct Models	7
4.6	Scoring Pipeline: Log-Probability Extraction	7
4.7	Quantization and Infrastructure	7
4.8	Evaluation Metrics	7
4.8.1	Expected Calibration Error (ECE)	7
4.8.2	Brier Score	8
4.8.3	C-index for Ordinal Validity	8
4.8.4	Reliability Diagrams	8
4.8.5	Bootstrap Confidence Intervals	8
4.9	Statistical Analysis Plan	8
4.9.1	Pre-registration of Hypotheses	8
4.9.2	Permutation Tests	8
4.9.3	Bonferroni Correction	8
4.9.4	Post-hoc Recalibration Methods	8
4.9.5	Conformal Prediction	8

4.10	Reproducibility	8
5	Results and Analysis	9
5.1	Chapter Overview	9
5.2	Pre-LLM Audit and Expert Agreement	9
5.3	Calibration Audit: H1 Supported Across All Models	9
5.4	The Effect of RLHF: H2 Supported (Main Contribution)	10
5.5	The Enriched-Prompt Paradox	10
5.6	Mode Collapse in Medical and RLHF-Aligned LLMs	10
5.7	Performance on Truly Discriminative Cases	10
5.8	Systematic Inclusion vs Exclusion Asymmetry	11
5.9	Model-Specific Evasion Biases	11
5.10	Ordinal Validity (H3) and Selective Classification (H4)	12
5.10.1	H3: Not Reached but Revealing Hierarchy	12
5.10.2	H4: Supported Only on the Mistral Family	12
5.11	Post-hoc Recalibration and Stability	12
5.12	Conformal Prediction and Inter-Model Consensus	12
5.12.1	Conformal Prediction Achieves Target Coverage	12
5.12.2	Inter-Model Consensus	12
5.13	Baselines and Composite Utility Score	12
5.14	Summary of Results	12
6	Conclusion and Future Work	13
6.1	Summary of Findings	13
6.2	Answer to the Research Question	13
6.3	Contributions	13
6.4	Limitations	13
6.4.1	External Validation	13
6.4.2	Model Scale	14
6.4.3	Prompt Space	14
6.5	Future Work	14
6.5.1	External Validation on Independent Benchmarks	14
6.5.2	Multi-Agent Adversarial Debate for Calibration	14
6.5.3	Investigating the Origin of Mode Collapse	14
6.5.4	Extension to Larger and Multilingual Models	14
6.6	Closing Remarks	14
	Bibliography	15
	A Detailed Results Tables	16

B	Full Prompt Specifications	17
B.1	Prompt A — Generic (Baseline)	17
B.2	Prompt B — Adapted with Expert Guidance	17
C	Individual Reliability Diagrams	18
D	Code and Data Availability	19
D.1	Repository Structure	19
D.2	Data Availability	19
D.3	Reproducing the Experiments	19
D.4	Computational Requirements	19
E	Generative AI Usage Declaration	20
F	Detailed Mode Collapse Analysis	22
F.1	Prediction Distribution per Configuration	22
F.2	Simulation: Forcing Decisive Predictions	22
F.3	Composite Utility Score Ranking	22
F.4	Cross-Validation Stability	22

List of Figures

- 4.1 Overview of the experimental design: 6 models $\times$ 2 prompts = 12 configurations evaluated on 1,015 patient-criterion pairs. 6
- 5.1 Reliability diagrams of all 12 configurations (Prompt A left, Prompt B right). All curves fall well below the perfect calibration diagonal. 9
- 5.2 Risk-coverage curves for the 6 Prompt-B configurations. 12
- C.1 Reliability diagrams for the 12 configurations. 18

List of Tables

2.1	Comparative positioning of related works and this thesis.	4
4.1	Selected models and their scientific role.	7
5.1	Calibration metrics across the 12 configurations.	9
5.2	Global vs. non-trivial accuracy across configurations. The performance drop reveals which models rely on the majority class distribution. . . .	11

List of Abbreviations

Acronym	Definition
AUC	Area Under the Curve
CI	Confidence Interval
ECE	Expected Calibration Error
EHR	Electronic Health Record
GPU	Graphics Processing Unit
IC	Intervalle de Confiance (French)
LLM	Large Language Model
LoRA	Low-Rank Adaptation
MC	Mode Collapse
ML4H	Machine Learning for Health
NEI	Not Enough Information
NF4	Normal Float 4-bit (quantization)
NLL	Negative Log-Likelihood
NLP	Natural Language Processing
PPO	Proximal Policy Optimization
QA	Question Answering
RCT	Randomized Controlled Trial
RLHF	Reinforcement Learning from Human Feedback
SFT	Supervised Fine-Tuning
TREC	Text REtrieval Conference

Chapter 1

Introduction

- 1.1 Macro Context: AI in Healthcare
- 1.2 Micro Context: Clinical Trial Matching
- 1.3 Problem Statement
- 1.4 Research Question and Hypotheses
- 1.5 Contributions
- 1.6 Structure of the Thesis

Chapter 2

State of the Art

2.1 Patient–Clinical Trial Matching

2.1.1 Clinical Context and Need

2.1.2 Pre-LLM Approaches

2.1.3 LLM-based Approaches

2.2 Probabilistic Calibration of Classifiers

2.2.1 Definitions and Metrics

2.2.2 Clinical Importance of Calibration

2.2.3 Recalibration Methods

2.3 Calibration of Large Language Models

2.3.1 The RLHF Effect on Calibration

2.3.2 Calibration of Medical LLMs

2.3.3 Identified Gap

2.4 Open-Source Medical LLMs

2.4.1 The Meditron Family

2.4.2 BioMistral and Model Merges

2.4.3 Other Open-Source Medical LLMs

2.5 Positioning of This Work

Table 2.1: Comparative positioning of related works and this thesis.

Work	Domain	Models	RLHF	Calib.	Recalib.
Jin et al. 2024	Trial match- ing	GPT-4	—	—	—
Kadavath 2022	QA	Anthropic	✓	✓	—
Labrak et al. 2024	Med QA	BioMistral	—	✓	—
This work	Trial match- ing	6 LLMs	✓	✓	✓

Chapter 3

Project Context and Scope

3.1 Academic Research Framework

3.2 Project Origin and Evolution

3.3 Dataset: TrialGPT-Criterion-Annotations

3.3.1 Dataset Description

3.3.2 Justification of the Single-Dataset Choice

3.4 Ethical Considerations

3.5 Constraints and Resources

3.5.1 Computational Constraints

3.5.2 Model Access Constraints

3.5.3 Time Constraints

3.6 Scope of This Work

Chapter 4

Methodology and Technical Approach

4.1 Overview of the Experimental Design

Figure 4.1: Overview of the experimental design: 6 models $\times$ 2 prompts = 12 configurations evaluated on 1,015 patient-criterion pairs.

4.2 Research Question and Pre-registered Hypotheses

Hypothesis 1 (H1). *Current LLM pipelines are poorly calibrated in clinical trial matching: $ECE > 0.05$ (threshold from Van Calster et al. 2019).*

Hypothesis 2 (H2). *RLHF post-training degrades calibration: $ECE(\text{Instruct}) > ECE(\text{Base})$, tested by permutation with Bonferroni correction.*

Hypothesis 3 (H3). *LLM eligibility scores exhibit ordinal validity: $C\text{-index} > 0.70$.*

Hypothesis 4 (H4). *Selective abstention improves clinical quality: $F_1(80\% \text{ coverage}) > F_1(100\% \text{ coverage}) + 0.02$.*

Table 4.1: Selected models and their scientific role.

Model	Type	Scientific Role
MISTRAL-7B-v0.3	Generalist, no RLHF	H2 control
MISTRAL-7B-INSTRUCT-v0.3	Generalist, with RLHF	H2 test
LLAMA-3.1-8B	Generalist, no RLHF	H2 replication
LLAMA-3.1-8B-INSTRUCT	Generalist, with RLHF	H2 replication
MEDITRON-7B	Medical, no RLHF	Pure specialization effect
BIO-MISTRAL-7B	Medical, inherited RLHF	Validation vs Labrak 2024

4.3 Data Preparation

4.4 Model Selection

4.4.1 Selection Criteria

4.4.2 The Six Selected Models

4.4.3 Excluded Models

4.5 Prompt Design

4.5.1 Rationale for Two Prompt Variants

4.5.2 Prompt A: Generic Baseline

4.5.3 Prompt B: Adapted with Expert Guidance

4.5.4 Adaptations for Base vs Instruct Models

4.6 Scoring Pipeline: Log-Probability Extraction

4.7 Quantization and Infrastructure

4.8 Evaluation Metrics

4.8.1 Expected Calibration Error (ECE)

$$\text{ECE} = \sum_{k=1}^K \frac{|B_k|}{n} |\text{acc}(B_k) - \text{conf}(B_k)| \quad (4.1)$$

Algorithm 1 Log-probability scoring for eligibility label prediction

Require: Patient note n , criterion c , criterion type t

Ensure: Predicted label $\hat{y}$, confidence $\hat{p}$, distribution π

```

1: Build prompt from  $n, c, t$ 
2:  $\mathcal{L} \leftarrow$  candidate labels according to  $t$ 
3: for all  $\ell \in \mathcal{L}$  do
4:    $s_\ell \leftarrow$  mean log-probability of tokens of  $\ell$ 
5: end for
6:  $\pi \leftarrow \text{softmax}(s)$ 
7:  $\hat{y} \leftarrow \arg \max_{\ell} \pi_\ell$ 
8:  $\hat{p} \leftarrow \pi_{\hat{y}}$ 
9: return  $\hat{y}, \hat{p}, \pi$ 

```

4.8.2 Brier Score

$$\text{Brier} = \frac{1}{n} \sum_{i=1}^n (p_i - y_i)^2 \quad (4.2)$$

4.8.3 C-index for Ordinal Validity

4.8.4 Reliability Diagrams

4.8.5 Bootstrap Confidence Intervals

4.9 Statistical Analysis Plan

4.9.1 Pre-registration of Hypotheses

4.9.2 Permutation Tests

4.9.3 Bonferroni Correction

4.9.4 Post-hoc Recalibration Methods

4.9.5 Conformal Prediction

4.10 Reproducibility

Chapter 5

Results and Analysis

5.1 Chapter Overview

5.2 Pre-LLM Audit and Expert Agreement

Finding 1 (F1: GPT-4’s Selective Evasion Bias). *GPT-4 over-uses the “not applicable” label (ratio 2.16 vs experts) but uses “not enough information” with a ratio of 0.99. This asymmetric pattern reveals a selective evasion, unreported in the original TrialGPT paper.*

5.3 Calibration Audit: H1 Supported Across All Models

Table 5.1: Calibration metrics across the 12 configurations.

Configuration	ECE	95% CI	Brier
Mistral-Base A	0.133	[0.104, 0.163]	0.258
Mistral-Instruct A	0.322	[0.297, 0.350]	0.311
Mistral-Base B	0.189	[0.161, 0.218]	0.269
Mistral-Instruct B	0.253	[0.229, 0.287]	0.353
Llama-Base A	0.124	[0.100, 0.148]	0.187
Llama-Instruct A	0.510	[0.483, 0.535]	0.456
Llama-Base B	0.151	[0.123, 0.177]	0.225
Llama-Instruct B	0.602	[0.575, 0.628]	0.558
Meditron A	0.182	[0.153, 0.212]	0.263
Meditron B	0.231	[0.204, 0.259]	0.283
BioMistral A	0.155	[0.126, 0.182]	0.234
BioMistral B	0.194	[0.166, 0.222]	0.244

Figure 5.1: Reliability diagrams of all 12 configurations (Prompt A left, Prompt B right). All curves fall well below the perfect calibration diagonal.

5.4 The Effect of RLHF: H2 Supported (Main Contribution)

Finding 2 (F2: RLHF Systematically Degrades Calibration). *Across two independent model families (Mistral-7B and Llama-3.1-8B) and two independent prompt designs, RLHF-aligned Instruct versions exhibit significantly worse calibration than their Base counterparts (all $p < 0.003$ with Bonferroni correction, ΔECE ranging from $+0.064$ to $+0.451$).*

5.5 The Enriched-Prompt Paradox

Finding 3 (F3: The Enriched-Prompt Paradox). *Contrary to the intuition that expert prompt design improves calibration, prompts enriched with definitions, examples, and exclusion rules degrade calibration in 5 of 6 models. Only the strongly RLHF-aligned Mistral-Instruct benefits.*

5.6 Mode Collapse in Medical and RLHF-Aligned LLMs

Finding 4 (F7: Mode Collapse in 5 of 12 Configurations). *Five open-source configurations — both BioMistral variants, all Llama-Instruct variants, and Llama-Base with prompt B — predict “not enough information” in more than 97% of cases. Their apparently reasonable ECE (0.15–0.60) reflects proximity to a trivial NEI baseline rather than genuine probabilistic quality.*

Finding 5 (F8: Mode Collapse Is Structural, Not Superficial). *When forced to choose between the two decisive labels of each criterion type by removing evasion candidates from the softmax, collapsed models do not recover reasoning ability. Instead, collapse migrates from “not enough information” to the majority label per criterion type (“not excluded” for exclusion criteria, or “excluded” for one Llama-Base variant). Evasion behaviors mask absence of clinical reasoning rather than strategic uncertainty communication.*

5.7 Performance on Truly Discriminative Cases

Finding 6 (F9: Global Accuracy Hides Complete Failure on Discriminative Cases). *On the 238 truly discriminative cases (23.4% of the dataset), Meditron achieves 0% accuracy despite its 64.8% global accuracy; BioMistral achieves 0.0% and 0.4% accuracy for its two prompts. The apparent superiority of medical LLMs disappears on the*

cases requiring genuine clinical decisions. Only Mistral-Instruct with prompt A reaches acceptable clinical performance (71.9% accuracy, ECE 0.054).

Table 5.2: Global vs. non-trivial accuracy across configurations. The performance drop reveals which models rely on the majority class distribution.

Configuration	Global acc.	Non-trivial acc.	Drop
Meditron A	0.648	0.000	−0.847
Meditron B	0.630	0.000	−0.824
BioMistral B	0.289	0.000	−0.377
BioMistral A	0.296	0.004	−0.381
Llama-Instruct B	0.292	0.013	−0.365
Llama-Base B	0.308	0.013	−0.386
Llama-Instruct A	0.290	0.021	−0.351
Llama-Base A	0.230	0.265	+0.046
Mistral-Base B	0.586	0.416	−0.222
Mistral-Instruct B	0.480	0.500	+0.026
Mistral-Base A	0.500	0.613	+0.149
Mistral-Instruct A	0.308	0.719	+0.536

5.8 Systematic Inclusion vs Exclusion Asymmetry

Finding 7 (F4: Inclusion vs Exclusion Asymmetry). *LLM calibration on exclusion criteria is 2–3 times worse than on inclusion criteria across all generalist models. MEDITRON is the sole exception, showing the reverse pattern — likely due to its clinical guidelines training data.*

5.9 Model-Specific Evasion Biases

Finding 8 (F5: Divergent Evasion Strategies Across LLMs). *Whereas GPT-4 evades via the “not applicable” label (ratio 2.16), open- source LLMs consistently evade via “not enough information” (ratio 3.3 to 3.5 for BioMistral and Llama-Instruct). Evasion is systematic but channel-specific.*

Figure 5.2: Risk-coverage curves for the 6 Prompt-B configurations.

5.10 Ordinal Validity (H3) and Selective Classification (H4)

5.10.1 H3: Not Reached but Revealing Hierarchy

5.10.2 H4: Supported Only on the Mistral Family

5.11 Post-hoc Recalibration and Stability

Finding 9 (F6: Post-hoc Isotonic Recalibration Works). *Isotonic regression fitted on 70% of the data brings 11 of 12 configurations below the well-calibrated threshold of 0.05, robustly across 5-fold cross-validation. However, for collapsed models the “correction” merely calibrates a near-constant prediction and does not restore classification quality.*

5.12 Conformal Prediction and Inter-Model Consensus

5.12.1 Conformal Prediction Achieves Target Coverage

5.12.2 Inter-Model Consensus

5.13 Baselines and Composite Utility Score

5.14 Summary of Results

Chapter 6

Conclusion and Future Work

6.1 Summary of Findings

6.2 Answer to the Research Question

6.3 Contributions

6.4 Limitations

6.4.1 External Validation

The present study relies on a single benchmark, TrialGPT-Criterion-Annotations [1], chosen because it is the only public corpus offering per-criterion multi-level expert annotations with an explicit inclusion/exclusion distinction — properties that are essential for fine-grained calibration analysis. External validation on independent datasets was not carried out, for three complementary reasons.

First, existing alternatives are not directly comparable. TREC Clinical Trials 2021/2022 and SIGIR 2016 operate at the trial level rather than the criterion level and use only three ordinal relevance grades, which precludes direct replication of the calibration protocol used here. Adapting the hypotheses to these datasets would not constitute a validation but a distinct study with different endpoints.

Second, the mode collapse phenomenon documented in Section 5.6 appears intrinsically linked to the availability of an evasion label (“not enough information”) in the label space. Datasets that do not include such a label cannot be used to test whether this behavior generalizes: they would trivially prevent it by design.

Third, prospective validation on real clinical data would require ethical approval, hospital partnership, and integration into a live recruitment workflow. These constraints situate external validation as a research project in its own right, complementary to the present work rather than a subordinate step.

External validation nevertheless remains an important direction for future work.

6.4.2 Model Scale

6.4.3 Prompt Space

6.5 Future Work

6.5.1 External Validation on Independent Benchmarks

6.5.2 Multi-Agent Adversarial Debate for Calibration

6.5.3 Investigating the Origin of Mode Collapse

6.5.4 Extension to Larger and Multilingual Models

6.6 Closing Remarks

Bibliography

- [1] Q. Jin, Z. Wang, et al., “Matching patients to clinical trials with large language models”, *Nature Communications*, vol. 15, p. 9074, 2024.

Appendix A

Detailed Results Tables

Appendix B

Full Prompt Specifications

B.1 Prompt A — Generic (Baseline)

```
1 Classify whether this patient meets the eligibility criterion.
2
3 Patient: {note}
4 Criterion ({type}): {criterion}
5
6 Choose exactly one label from: included, not included, excluded
7   ,
8   not excluded, not enough information, not applicable.
9
10 Answer :
```

Listing B.1: Prompt A for Instruct models

B.2 Prompt B — Adapted with Expert Guidance

Appendix C

Individual Reliability Diagrams

(a) Mistral-Base, Prompt A

(b) Mistral-Base, Prompt B

Figure C.1: Reliability diagrams for the 12 configurations.

Appendix D

Code and Data Availability

D.1 Repository Structure

D.2 Data Availability

D.3 Reproducing the Experiments

D.4 Computational Requirements

Appendix E

Generative AI Usage Declaration

Tools Used

During the preparation of this thesis, the following generative AI tool was used as an assistant:

- **Claude** (Anthropic), primarily the Sonnet and Opus series model families, accessed via the `claude.ai` interface.

Purpose of Use

Generative AI was used for:

- **Ideation and design discussions:** brainstorming experimental designs, discussing methodological trade-offs, and refining the research question.
- **Code assistance:** writing and debugging Python code for the evaluation pipeline, prompt engineering iterations, and statistical analyses.
- **Writing assistance:** refining the phrasing of prompts, drafting progress reports, and revising English formulations of methodological descriptions.
- **Bibliographic exploration:** identifying candidate references and papers relevant to the state of the art.

Purpose of Non-Use

Generative AI was *not* used for:

- Generating experimental data or results.
- Fabricating statistical outcomes or numerical values.
- Producing scientific claims not backed by our own experiments.

- Writing the final thesis text verbatim (the author remains responsible for all formulations).

Verification and Author Responsibility

All AI-generated content used in this work was:

- Reviewed, verified, and adapted by the author before inclusion.
- Cross-checked against primary sources for any factual claim.
- Understood by the author, who takes full responsibility for the scientific and methodological choices reflected in this thesis.

The scientific reasoning, experimental design, execution, and interpretation of results are the sole responsibility of the author.

Appendix F

Detailed Mode Collapse Analysis

F.1 Prediction Distribution per Configuration

F.2 Simulation: Forcing Decisive Predictions

F.3 Composite Utility Score Ranking

F.4 Cross-Validation Stability

EOF echo "OK"